

超大卷积核架构设计与高效实践

MEGVII 旷视

——解读Scaling Up Your Kernels to **31x31**:
Revisiting Large Kernel Design in CNNs
(CVPR 2022)

主讲人：丁霄汉

目录

Contents

- 1 **WHY** : 关于CNN和Transformer的基本问题
- 2 **HOW** : 大kernel设计探索与五大准则
- 3 RepLKNet
- 4 实验结果
- 5 讨论与分析
- 6 总结

/ 01

“ 关于CNN和 Transformer的基本问题 ”

问题一：与CNN相比，Transformer性能强的原因是什么？

Transformer的基本组件是self-attention，而self-attention的实质是在**全局尺度或较大的窗口内**进行**Query-Key-Value运算**。其中的关键是？

A. 全局尺度或较大的窗口

B. Query-Key-Value运算

证据一：将Swin中的attention换成**7x7卷积**，性能也很强。

证据二：将ViT中的attention换成**MLP**，性能也很强。

证据三：将attention换成**pooling**，性能还是很强。

[1] Qi Han, Zejia Fan, Qi Dai, Lei Sun, Ming-Ming Cheng, Jiaying Liu, and Jingdong Wang. Demystifying local vision transformer: Sparse connectivity, weight sharing, and dynamic weight.

[2] Ilya Tolstikhin, Neil Houlsby, Alexander Kolesnikov, Lucas Beyer, Xiaohua Zhai, Thomas Unterthiner, Jessica Yung, Andreas Steiner, Daniel Keysers, Jakob Uszkoreit, et al. Mlp-mixer: An all-mlp architecture for vision.

[3] Weihao Yu, Mi Luo, Pan Zhou, Chenyang Si, Yichen Zhou, Xinchao Wang, Jiashi Feng, and Shuicheng Yan. Metaformer is actually what you need for vision.

问题二：如果感受野是关键，那么传统CNN差在哪里了？

深层CNN的感受野不是很大吗？

自VGGNet以来，堆叠小kernel成为了主流设计范式

- AlexNet : 11x11
- VGGNet : 3x3
- ResNet、ResNeXt : 一层7x7，大量3x3
- EfficientNet : 3x3、5x5
- MobileNet、ShuffleNet : 3x3
- RegNet : 3x3


问题二：如果感受野是关键，那么传统CNN差在哪里了？

因为我们相信，三层3x3比一层7x7要好：

- $3 \times 3 \times 3 < 1 \times 7 \times 7$
- 三层3x3可以加入更多非线性，层数越多，拟合能力越强
- 两者的感受野一样大（？）
- 在ResNet等深度现代模型中，大量3x3堆出来的感受野是足够覆盖全图的（？）

这是真的吗？

问题二：如果感受野是关键，那么传统CNN差在哪里了？



有效感受野可视化

小kernel模型增加深度无法显著增大有效感受野，大kernel模型的有效感受野非常大

有效感受野正比于 $K\sqrt{L}$ ，K为kernel size，L为深度

问题三：为什么我们不用大kernel？

为什么历史淘汰了大kernel？

- **太大**：卷积的参数量和FLOPs与kernel size的平方成正比，如(256, 256, 31, 31)的一个卷积核就有63M参数
- **反而掉点**：
 - [1] GCN：Kx1和1xK大卷积，ImageNet上掉点
 - [2] LR-Net：大kernel，非卷积，kernel size从7x7到9x9即开始掉点

[1] Chao Peng, Xiangyu Zhang, Gang Yu, Guiming Luo, and Jian Sun. Large kernel matters—improve semantic segmentation by global convolutional network. CVPR 2017

[2] Han Hu, Zheng Zhang, Zhenda Xie, and Stephen Lin. Local relation networks for image recognition. CVPR 2019

问题三：为什么我们不用大kernel？

曾被历史淘汰 \neq 不应该得到复兴

时代变了：

- 太大：
 - Depth-wise卷积
 - 恰当的底层优化
- 反而掉点：
 - 配套的结构设计
 - 重参数化等新方法

问题四：如何在现代模型中复兴大kernel设计？

大kernel不是不能用，而是需要用到好处

How？

/ 02

“大kernel设计探索与 五大准则”

准则一：depth-wise卷积+恰当的底层优化

- 从3x3改到31x31，模型会增大100倍？
- Depth-wise (DW) 卷积占的参数量和FLOPs本来就不多！

Table 5. RepLKNet with different kernel sizes. The models are pretrained on ImageNet-1K in 120 epochs with 224×224 input and finetuned on ADE20K with UperNet in 80K iterations. On ADE20K, we test the *single-scale* mIoU, and compute the FLOPs with input of 2048×512 , following Swin.

Kernel size	ImageNet			ADE20K		
	Top-1	Params	FLOPs	mIoU	Params	FLOPs
3-3-3-3	82.11	71.8M	12.9G	46.05	104.1M	1119G
7-7-7-7	82.73	72.2M	13.1G	48.05	104.6M	1123G
13-13-13-13	83.02	73.7M	13.4G	48.35	106.0M	1130G
25-25-25-13	83.00	78.2M	14.8G	48.68	110.6M	1159G
31-29-27-13	83.07	79.3M	15.3G	49.17	111.7M	1170G

准则一：depth-wise卷积+恰当的底层优化

- 加底层优化，速度更快
 - Depth-wise卷积效率低？DW 3x3效率低，不代表DW 31x31效率低。
 - 优化得当，可加速20倍

Table 1. Inference speed of a stack of 24-layer depth-wise convolutions with various kernel sizes and resolutions on a single GTX 2080Ti GPU. The input shape is (64, 384, R , R). Baselines are evaluated with Pytorch 1.9.0 + cuDNN 7.6.5, in FP32 precision.

Resolution R	Impl	Latency (ms) @ Kernel size									
		3	5	7	9	13	17	21	27	29	31
16×16	Pytorch	5.6	11.0	14.4	17.6	36.0	57.2	83.4	133.5	150.7	171.4
	Ours	5.6	6.5	6.4	6.9	7.5	8.4	8.4	8.4	8.3	8.4
32×32	Pytorch	21.9	34.1	54.8	76.1	141.2	230.5	342.3	557.8	638.6	734.8
	Ours	21.9	28.7	34.6	40.6	52.5	64.5	73.9	87.9	92.7	96.7
64×64	Pytorch	69.6	141.2	228.6	319.8	600.0	977.7	1454.4	2371.1	2698.4	3090.4
	Ours	69.6	112.6	130.7	152.6	199.7	251.5	301.0	378.2	406.0	431.7

- 已集成进开源框架MegEngine
- 开源的PyTorch实现，参见示例：<https://github.com/DingXiaoH/RepLKNet-pytorch>



准则二：加shortcut！

- 实验：MobileNet V2，有/无shortcut，加大到13x13

Table 2. Results of different kernel sizes in normal/shortcut-free MobileNet V2.

Shortcut	Kernel size	ImageNet top-1 accuracy (%)
✓	3×3	71.76
✓	13×13	72.53
	3×3	68.67
	13×13	53.98

- 解释：shortcut产生了“**组合式**”的感受野；没有shortcut，感受野单一且巨大，难以捕捉小的特征

准则三：用小kernel做重参数化

- 实验：MobileNet V2，有/无3x3重参数化，加大到9x9或13x13
- 重参数化：
 - 训练时有并行的小kernel，如3x3
 - 训练后等效地将小kernel叠加到大kernel的中心
 - 原理：卷积的可加性




Figure 2. An example of re-parameterizing a small kernel (e.g., 3×3) into a large one (e.g., 7×7). See [27, 30] for details.

准则三：用小kernel做重参数化

- 效果：
 - 无重参数化：9x9涨点，13x13掉点
 - 有重参数化：**13x13涨点**
- 解释：小kernel使组合式的感受野更加丰富，更容易提取小的特征
- 注：在数据量很大 (MegData-73M) 时，小kernel重参数化 (inductive bias) 效果不大

Table 3. Results of 3×3 re-parameterization on MobileNet V2 with various kernel sizes.

Kernel	3×3 re-param	ImageNet top-1 acc (%)	Cityscapes val mIoU (%)
3×3	N/A	71.76	72.31
9×9		72.67	76.11
9×9	✓	73.09	76.30
13×13		72.53	75.67
13×13	✓	73.24	76.60

准则四：看下游任务！

- 解释：
 - ImageNet分类需要的信息量不大，更大的感受野和更强的表征能力不一定能转化为涨点，但下游需要的语义信息更加丰富。
 - ImageNet模型主要靠纹理来做决定

Table 3. Results of 3×3 re-parameterization on MobileNet V2 with various kernel sizes.

Kernel	3×3 re-param	ImageNet top-1 acc (%)	Cityscapes val mIoU (%)
3×3	N/A	71.76	72.31
9×9		72.67	76.11
9×9	✓	73.09	76.30
13×13		72.53	75.67
13×13	✓	73.24	76.60

Table 5. RepLKNet with different kernel sizes. The models are pretrained on ImageNet-1K in 120 epochs with 224×224 input and finetuned on ADE20K with UperNet in 80K iterations. On ADE20K, we test the *single-scale* mIoU, and compute the FLOPs with input of 2048×512, following Swin.

Kernel size	ImageNet			ADE20K		
	Top-1	Params	FLOPs	mIoU	Params	FLOPs
3-3-3-3	82.11	71.8M	12.9G	46.05	104.1M	1119G
7-7-7-7	82.73	72.2M	13.1G	48.05	104.6M	1123G
13-13-13-13	83.02	73.7M	13.4G	48.35	106.0M	1130G
25-25-25-13	83.00	78.2M	14.8G	48.68	110.6M	1159G
31-29-27-13	83.07	79.3M	15.3G	49.17	111.7M	1170G

Robert Geirhos, Patricia Rubisch, Claudio Michaelis, Matthias Bethge, Felix A Wichmann, and Wieland Brendel.
Imagenet-trained cnns are biased towards texture; increasing shape bias improves accuracy and robustness.

准则五：小feature map上也能用大kernel！

- 大kernel一定要用大分辨率？不需要！
- 7x7的feature map上用13x13的大kernel？不严格符合平移不变性
 - 平移不变性一定是好的？
 - 视角一：可以看成一种不同位置卷积核参数不同的特殊卷积
 - 视角二：padding引入了绝对位置信息




Figure 3. Illustration to convolution with small feature map and large kernel. Two outputs at adjacent locations only share a part of kernel weights. Translational equivariance does not strictly hold.

Table 4. Results of various kernel sizes in the *last stage* of MobileNet V2. Kernel sizes in previous stages remain to be 3×3 .

Kernel size	ImageNet acc (%)	Cityscapes mIoU (%)
3×3	71.76	72.31
7×7	72.00	74.30
13×13	71.97	74.62

准则一：depth-wise卷积+恰当的底层优化

准则二：加shortcut！

准则三：用小kernel做重参数化

准则四：看下游任务！

准则五：小feature map上也能用大kernel！

/ 03

“

RepLKNet

”

- 基于以上准则，设计RepLKNet
- 借鉴Swin的宏观架构
- CNN风格的一点改动

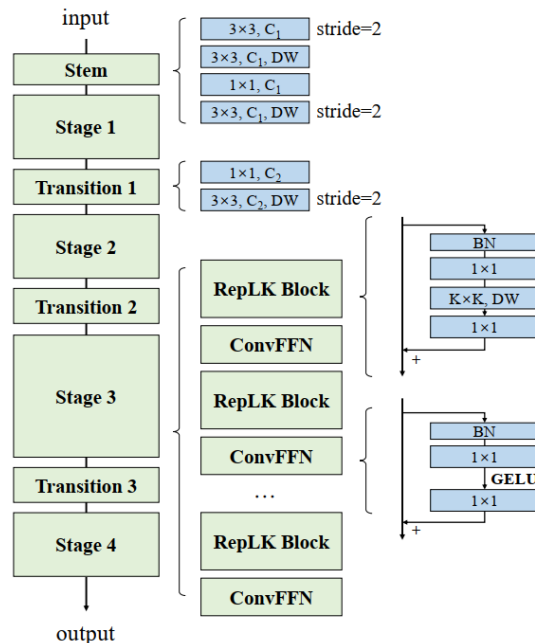


Figure 4. RepLKNet comprises Stem, Stages and Transitions. Except for depth-wise (DW) large kernel, the other components include DW 3×3 , dense 1×1 conv, and batch normalization [50] (BN). Note that every conv layer has a following BN, which are not depicted. Such conv-BN sequences use ReLU as the activation function, except those before the shortcut-addition (as a common practice [41, 73]) and those preceding GELU [42].

- 加大kernel size!
- RepLKNet-31B : 宽度[128, 256, 512, 1024]
- RepLKNet-31L : 宽度[192, 384, 768, 1536]
- RepLKNet-XL : 宽度[256, 512, 1024, 2048] + 1.5x inverted bottleneck

Table 5. RepLKNet with different kernel sizes. The models are pretrained on ImageNet-1K in 120 epochs with 224×224 input and finetuned on ADE20K with UperNet in 80K iterations. On ADE20K, we test the *single-scale* mIoU, and compute the FLOPs with input of 2048×512 , following Swin.

Kernel size	ImageNet			ADE20K		
	Top-1	Params	FLOPs	mIoU	Params	FLOPs
3-3-3-3	82.11	71.8M	12.9G	46.05	104.1M	1119G
7-7-7-7	82.73	72.2M	13.1G	48.05	104.6M	1123G
13-13-13-13	83.02	73.7M	13.4G	48.35	106.0M	1130G
25-25-25-13	83.00	78.2M	14.8G	48.68	110.6M	1159G
31-29-27-13	83.07	79.3M	15.3G	49.17	111.7M	1170G

/ 04

“

实验结果

”

- ImageNet

Table 6. ImageNet results. The throughput is tested with FP32 and a batch size of 64 on 2080Ti. ‡ indicates ImageNet-22K pretraining. ◇ indicates pretrained with extra data.

Model	Input resolution	Top-1 acc	Params (M)	FLOPs (G)	Throughput examples/s
RepLKNet-31B	224×224	83.5	79	15.3	295.5
Swin-B	224×224	83.5	88	15.4	226.2
RepLKNet-31B	384×384	84.8	79	45.1	97.0
Swin-B	384×384	84.5	88	47.0	67.9
RepLKNet-31B ‡	224×224	85.2	-	-	-
Swin-B ‡	224×224	85.2	-	-	-
RepLKNet-31B ‡	384×384	86.0	-	-	-
Swin-B ‡	384×384	86.4	-	-	-
RepLKNet-31L ‡	384×384	86.6	172	96.0	50.2
Swin-L ‡	384×384	87.3	197	103.9	36.2
RepLKNet-XL ◇	320×320	87.8	335	128.7	39.1

- Cityscapes
 - RepLKNet-31**Base** + ImageNet-**1K** > Swin-**Large** + ImageNet-**22K**

Table 7. Cityscapes results. The FLOPs is computed with 1024×2048 inputs. The mIoU is tested with single-scale (ss) and multi-scale (ms). The results with Swin are implemented by [37].

‡ indicates ImageNet-22K pretraining.

Backbone	Method	mIoU (ss)	mIoU (ms)	Param (M)	FLOPs (G)
RepLKNet-31B	UperNet [98]	83.1	83.5	110	2315
ResNeSt-200	DeepLabv3	-	82.7	-	-
Axial-Res-XL	Axial-DL [91]	80.6	81.1	173	2446
Swin-B	UperNet	80.4	81.5	121	2613
Swin-B	UperNet + [37]	80.8	81.8	121	-
ViT-L ‡	SETR-PUP [112]	79.3	82.1	318	-
ViT-L ‡	SETR-MLA	77.2	-	310	-
Swin-L ‡	UperNet	82.3	83.1	234	3771
Swin-L ‡	UperNet + [37]	82.7	83.6	234	-

- ADE20K
 - 仅ImageNet-1K pretraining, **50.6 mIoU**
 - ViT-L级别 + 73M数据, **56.0 mIoU**

Table 8. ADE20K results. The mIoU is tested with single-scale (ss) and multi-scale (ms). The results with 1K-pretrained Swin are cited from the official GitHub repository. ‡ indicates ImageNet-22K pretraining and 640×640 finetuning on ADE20K. ◊ indicates pretrained with extra data. The FLOPs is computed with 2048×512 for the ImageNet-1K pretrained models and 2560×640 for the ImageNet-22K and larger, following Swin.

Backbone	Method	mIoU (ss)	mIoU (ms)	Param (M)	FLOPs (G)
RepLKNet-31B	UperNet	49.9	50.6	112	1170
ResNet-101	UperNet [98]	43.8	44.9	86	1029
ResNeSt-200	DeepLabv3	-	48.4	113	1752
Swin-B	UperNet	48.1	49.7	121	1188
Swin-B	UperNet + [37]	48.4	50.1	121	-
ViT-Hybrid	DPT-Hybrid	-	49.0	90	-
ViT-L	DPT-Large	-	47.6	307	-
ViT-B	SETR-PUP [112]	46.3	47.3	97	-
ViT-B	SETR-MLA [112]	46.2	47.7	92	-
RepLKNet-31B ‡	UperNet	51.5	52.3	112	1829
Swin-B ‡	UperNet	50.0	51.6	121	1841
RepLKNet-31L ‡	UperNet	52.4	52.7	207	2404
Swin-L ‡	UperNet	52.1	53.5	234	2468
ViT-L ‡	SETR-PUP	48.6	50.1	318	-
ViT-L ‡	SETR-MLA	48.6	50.3	310	-
RepLKNet-XL ◊	UperNet	55.2	56.0	374	3431

- COCO
 - 与Swin相当
 - 与ResNeXt-101-64x4d相比，体量更小，**AP高4.4**

Table 9. Object detection on COCO. The FLOPs is computed with 1280×800 inputs. The FCOS model is trained with the 2x (24-epoch) training schedule for a fair comparison with the X101 (short for ResNeXt-101) baseline from the same code base [19], and the other results with Cascade Mask R-CNN all use 3x (36-epoch). The results of X101-64x4d + Cas Mask are reported by [60]. The results of 22K-pretrained Swin (without HTC++ [60]) are reported by [61]. ‡ indicates ImageNet-22K pretraining. ◇ indicates pretrained with extra data.

Backbone	Method	AP ^{box}	AP ^{mask}	Param (M)	FLOPs (G)
RepLKNet-31B	FCOS	47.0	-	87	437
X101-64x4d	FCOS	42.6	-	90	439
RepLKNet-31B	Cas Mask	52.2	45.2	137	965
X101-64x4d	Cas Mask	48.3	41.7	140	972
ResNeSt-200	Cas R-CNN [9]	49.0	-	-	-
Swin-B	Cas Mask	51.9	45.0	145	982
RepLKNet-31B ‡	Cas Mask	53.0	46.0	137	965
Swin-B ‡	Cas Mask	53.0	45.8	145	982
RepLKNet-31L ‡	Cas Mask	53.9	46.5	229	1321
Swin-L ‡	Cas Mask	53.9	46.7	254	1382
RepLKNet-XL ◇	Cas Mask	55.5	48.0	392	1958

/ 05

“

讨论与分析

”

- 有效感受野
 - 定量分析：已经求出了每个像素的贡献值，那么，中间多大的一个区域包含了99%的贡献值？
 - ResNet-101 -> ResNet-152，增加约一半体量
 - RepLKNet-13 -> RepLKNet-31，增加约**7.6%**参数




Table 10. Quantitative analysis on the ERF with the high-contribution area ratio r . A larger r suggests a smoother distribution of high-contribution pixels, hence larger ERF.

	$t = 20\%$	$t = 30\%$	$t = 50\%$	$t = 99\%$
ResNet-101	1.0%	1.7%	3.5%	23.4%
ResNet-152	1.3%	2.1%	4.5%	34.9%
RepLKNet-13	10.2%	15.8%	28.5%	96.3%
RepLKNet-31	15.0%	23.4%	41.4%	98.6%

- Shape bias : 模型有多少决定是根据形状 (而非纹理) 做出的 ?
- Shape bias越高 , 跟人类越像
- 发现 :
 - Shape bias跟训练数据关系很大
 - RepLKNet-3和ResNet-152的shape bias几乎一样
 - **Swin的shape bias不高**
 - **RepLKNet-31的shape bias高**
- 已知ViT的shape bias高 (参见其论文)
- Shape bias跟感受野关系很大 ? 跟attention关系不大 ?




Figure 5. Shape bias of RepLKNet, Swin, and ResNet-152 pre-trained on ImageNet-1K or 22K. The scatters represent the shape bias of 16 categories, and the vertical lines are the averages across categories (note RepLKNet-3 and ResNet-152 are very close).

论文 : Shikhar Tuli, Ishita Dasgupta, Erin Grant, and Thomas L Griffiths. Are convolutional neural networks or transformers more like human vision?

开源工具 : <https://github.com/bethgelab/model-vs-human>

/ 06

“

总结

”

- 从小kernel到大kernel
 - Dense conv ? 太大了！↓用depthwise →还是用3x3吧
 - Depth-wise大kernel ? Depth-wise 3x3都那么慢！↓depth-wise大kernel不一定慢 →还是用3x3吧
 - 增大kernel size , 掉点了？↓加shortcut →还是用3x3吧
 - 加了shortcut , 从大kernel到超大kernel , 又掉点了？↓加重参数化/大数据 →还是用3x3吧
 - 大kernel肯定要大分辨率啊，训不动！↓正常分辨率直接开整 →还是用3x3吧
 - ImageNet还是不涨点？↓再看看下游的 →还是用3x3吧



<https://www.zhihu.com/question/517340666/answer/2390516334>

论文：<https://arxiv.org/abs/2203.06717>

代码 (MegEngine)：<https://github.com/megvii-research/RepLKNet>

代码 (PyTorch)：<https://github.com/DingXiaoH/RepLKNet-pytorch>

用人工智能造福大众

MEGVII 旷视